



# High-Performance Data-Intensive Computing Systems Laboratory

## With Big Data Comes Big Compute: Scaling Machine Learning onto Public and Commercial Clouds with Kubernetes - Part 2

Introduction to Containers and Kubernetes

IEEE International Conference on Big Data



University of Missouri

December 2023



# Part 2 - Learning Objectives

- ▶ Deploying Scikit Learn ML Jobs to Kubernetes
- ▶ Deploying GPU Jobs to Kubernetes for Training Computer Vision Models
- ▶ Scaling Computer Vision Models on Kubernetes with Job Automation





# MU-HPDI/Nautilus

- ▶ Sample Dockerfiles
- ▶ Sample Kubernetes YAML File
- ▶ Wiki with detailed walkthroughs for:
  - ▶ Getting Started
  - ▶ Creating PVC
  - ▶ Creating Pods
  - ▶ Creating Jobs
- ▶ Nautilus Portal:  
<https://portal.nrp-nautilus.io>
- ▶ JupyterHub Instance:  
<https://bigdata-2023.nrp-nautilus.io/>
- ▶ Tutorial Repository of Jupyter Project Pages, Code Samples, YAML, etc.  
<https://github.com/MU-HPDI/bigdata-2023>
- ▶ Git Clone Command:  
git clone <https://github.com/MU-HPDI/bigdata-2023>.git

**Over a short break, we will ensure everyone has cloned this Repo into their JupyterLab environment**







# High-Performance Data-Intensive Computing Systems Laboratory

Part 2-A

## National Research Platform Nautilus Research Cluster

Introduction to Containers and Kubernetes

IEEE International Conference on Big Data

December 2023



University of Missouri



# Motivation

- ▶ Shallow ML modules often require extensive hyperparameter optimization to find optimal performance
- ▶ Deep Learning models require incredible amounts of compute to effectively train
- ▶ Using single developer machines or local on-prem resources often fail to scale effectively
- ▶ Kubernetes Clusters provide an efficient and scalable solution to training ML and deep learning models at scale
- ▶ This workshop will address the building and deployment of ML and deep learning containers to Nautilus to train deep networks

# Part-2 Tutorial Outline

- ▶ Quick Containerization Review and Advanced:
  - ▶ Review of containerization
  - ▶ Building ML containers for Scikit-Learn
  - ▶ Building Optimized Containers for Deep Learning with PyTorch
  - ▶ Using Common Frameworks: Detectron2, MMDetection, and Ultralytics
- ▶ Kubernetes ML / Deep Learning:
  - ▶ Review of Kubernetes Architecture and Key Concepts
  - ▶ Introduction to NRP Nautilus HyperCluster
  - ▶ Migrating Data to NRP with S3
  - ▶ Deploying Scikit-Learn ML Jobs to Nautilus
  - ▶ Deploying GPU Jobs to Nautilus for Training Computer Vision Models





# Containerization

Building Optimized Containers for Deep Learning with PyTorch



# Containerization


## Docker

### ► Docker

- You can think of Docker containers as mini-VMs that contain all the packages, both at the OS and language-specific level, necessary to run your software.

### ► Nautilus Gitlab

- Offers the ability to create repo for your code
- Offers the ability to create easily maintained and developed custom images using CI/CD feature





# Using Common Deep Learning Frameworks


Open MMLab, FAIR, and Ultralytics



University of Missouri

# Deep Learning Frameworks

- ▶ Many companies and research labs are developing open source Deep Neural Network frameworks to perform Computer Vision tasks
  - ▶ Open MMLab
  - ▶ Facebook AI Research
  - ▶ Ultralytics
  - ▶ Google
- ▶ We can build optimized Docker containers with these frameworks and use them as a baseline for research
- ▶ To use a given framework on Nautilus, we only need a Dockerfile → From the Dockerfile, we can build a container image that can be deployed on the cluster





# MMLab Frameworks: MMDetection, MMClassification, MMSegmentation

- ▶ Open MMLab has developed a set of extensible libraries for performing key Computer Vision tasks:
  - ▶ Classification: MMClassification
  - ▶ Object Detection: MMDetection
  - ▶ Semantic Segmentation: MMSegmentation
- ▶ These libraries can serve as an excellent starting point for many CV applications
- ▶ Very large model zoo with community trained models and benchmarks
- ▶ Highly extensible and configurable frameworks



# MMDetection Dockerfile

- ▶ Dockerfile contains all steps to create a fully functional MMDetection Python environment
- ▶ Add your code to the container or use this Dockerfile as a base in a multi-stage build

```
1 FROM nvcr.io/nvidia/pytorch:22.06-py3
2
3 # arg
4 ARG MMDET_VERSION=2.25.0
5 ARG MMCV_VERSION=1.5.2
6 ARG MMCLS_VERSION=0.23.1
7 ARG BUILD_DIR="/build"
8
9 # env variables
10 ENV MPLCONFIGDIR /tmp
11 ENV TORCH_HOME /tmp
12 ENV MMCV_WITH_OPS 1
13 ENV FORCE_CUDA 1
14 ENV CUDA_HOME /usr/local/cuda
15
16 # create the build dir
17 RUN mkdir -p ${BUILD_DIR}
18
19 RUN apt update -y && apt install -y libpng-dev libjpeg-dev libgl-dev wget && rm -rf /var/lib/apt/lists/*
20 # install mm cv
21 # https://mmdcv.readthedocs.io/en/latest/get_started/build.html#build-on-linux-or-macos
22 RUN wget https://github.com/open-mmlab/mmcv/archive/refs/tags/v${MMCV_VERSION}.tar.gz -O /tmp/mmcv.tar.gz
23 RUN tar -xzf /tmp/mmcv.tar.gz
24 RUN mv ./mmcv-${MMCV_VERSION} ${BUILD_DIR}/mmcv
25 # install it
26 RUN pip install -r ${BUILD_DIR}/mmcv/requirements/optional.txt
27 RUN pip install -v ${BUILD_DIR}/mmcv
28
29
30 # install mm detection
31 # https://github.com/open-mmlab/mmdetection/blob/v2.25.0/docs/en/get_started.md#customize-installation
32 RUN wget https://github.com/open-mmlab/mmdetection/archive/refs/tags/v${MMDET_VERSION}.tar.gz -O /tmp/mmdet.tar.gz
33 RUN tar -xzf /tmp/mmdet.tar.gz
34 RUN mv ./mmdetection-${MMDET_VERSION} ${BUILD_DIR}/mmdet
35 # install it
36 RUN pip install ${BUILD_DIR}/mmdet
```



MU-HPDI/Nautilus



University of Missouri





- 



# Detectron2 Dockerfile

- ▶ Dockerfile contains all steps to create a fully functional Detectron2 Python environment
- ▶ Add your code to the container or use this Dockerfile as a base in a multi-stage build



```
1 FROM nvcr.io/nvidia/pytorch:21.06-py3
2
3 # install system reqs
4 RUN apt update && apt install -y vim libgl-dev
5 RUN apt-get install --reinstall ca-certificates # for git
6
7 # env variables
8 ENV MPLCONFIGDIR /tmp
9 ENV TORCH_HOME /tmp
10 ENV FVCORE_CACHE /tmp
11
12
13 #####
14 # Detectron 2
15 #####
16 RUN git clone -b 'v0.4' --single-branch --depth 1 --recursive https://github.com/facebookresearch/detectron2.git /workspace/detectron2
17 RUN pip install -v /workspace/detectron2/
```





# Ultralytics

- ▶ Developer of most used Open-Source implementation of YOLOv3, YOLOv5, and more recently, YOLOv8
- ▶ You Only Look Once (YOLO) is a common real-time object detection architecture with many variations over the years
  - ▶ YOLO
  - ▶ YOLO9000 (a.k.a. YOLOv2)
  - ▶ YOLOv3
  - ▶ YOLOv5
  - ▶ YOLOv8



# Ultralytics Dockerfile

- ▶ Dockerfile contains all steps to create a fully functional YOLOv3, YOLOv5, or YOLOv8 Python environment
- ▶ Add your code to the container or use this Dockerfile as a base in a multi-stage build



ultralytics/ultralytics



University of Missouri

```
1 # Ultralytics YOLO 🚀, AGPL-3.0 license
2 # Builds ultralytics/ultralytics:latest image on DockerHub https://hub.docker.com/r/ultralytics/ultralytics
3 # Image is CUDA-optimized for YOLOv8 single/multi-GPU training and inference
4
5 # Start FROM PyTorch image https://hub.docker.com/r/pytorch/pytorch or nvcr.io/nvidia/pytorch:23.03-py3
6 FROM pytorch/pytorch:2.0.0-cuda11.7-cudnn8-runtime
7
8 # Downloads to user config dir
9 ADD https://ultralytics.com/assets/Arial.ttf https://ultralytics.com/assets/Arial.Unicode.ttf /root/.config/Ultralytics/
10
11 # Install linux packages
12 # g++ required to build 'tflite_support' package
13 RUN apt update \
14     && apt install --no-install-recommends -y gcc git zip curl http libgl1-mesa-glx libglib2.0-0 libpython3-dev gnupg g++
15 # RUN alias python=python3
16
17 # Security updates
18 # https://security.snyk.io/vuln/SNYK-UBUNTU1804-OPENSSL-3314796
19 RUN apt upgrade --no-install-recommends -y openssl tar
20
21 # Create working directory
22 RUN mkdir -p /usr/src/ultralytics
23 WORKDIR /usr/src/ultralytics
24
25 # Copy contents
26 # COPY . /usr/src/app (issues as not a .git directory)
27 RUN git clone https://github.com/ultralytics/ultralytics /usr/src/ultralytics
28 ADD https://github.com/ultralytics/assets/releases/download/v0.0.0/yolov8n.pt /usr/src/ultralytics/
29
30 # Install pip packages
31 RUN python3 -m pip install --upgrade pip wheel
32 RUN pip install --no-cache -e . albumentations comet tensorboard thop pycocotools
33
34 # Set environment variables
35 ENV OMP_NUM_THREADS=1
```



# Migrating Data to Nautilus



# Deep Learning & Data

- ▶ Deep Learning algorithms require a vast amount of representative training data to effectively train
  - ▶ The more complex the network architecture, the more quality data is required
- ▶ Previous portions of this workshop have covered creating the containers to train the algorithms, but we have not yet covered how to stage data to Nautilus
  - ▶ All data for use on Nautilus will need to be moved to *persistent volumes* on the cluster
- ▶ Three ways to move data to a persistent volume on Nautilus:
  - ▶ Using KubeCTL
  - ▶ Using Commercial Cloud Storage
  - ▶ Using Nautilus S3 Storage ← *Recommended*





# Migrating Data to Nautilus: KubeCTL

- ▶ The command line Kubernetes tool, KubeCTL, has functionality to copy data to and from running pods
  - ▶ We can use this copy utility to move data to the cluster
- ▶ Advantages:
  - ▶ No additional installation or setup
  - ▶ No need to stage data in cloud storage
  - ▶ Requires only the KubeCTL command line tool
- ▶ Disadvantages:
  - ▶ Only small amounts (< 100 MB) of data can be moved per kubectl copy command
  - ▶ KubeCTL copy is *slow* at < 100 Mbps upload speeds

- ▶ How to:

```
kubectl cp localpath podname:/path/on/pod
```



# Migrating Data to Nautilus: Using Commercial Cloud Storage

- ▶ We can use commercial cloud storage, such GCP Buckets or AWS S3 to move data to Nautilus
- ▶ Advantages:
  - ▶ Flexibility of cloud platform
  - ▶ Very fast transfer speeds
  - ▶ Capable of virtually any amount of data
- ▶ Disadvantages
  - ▶ Cost
  - ▶ Setup of Cloud Storage
  - ▶ Installation of Cloud Interface
  - ▶ Staging of data in cloud
- ▶ How to:
  - ▶ Create cloud bucket with Commercial Cloud Vendor (i.e., GCP)
  - ▶ Copy data to Bucket: `gsutil cp localpath gs://bucketName`
  - ▶ Create Google Cloud Pod on Nautilus
  - ▶ Copy data from Bucket: `gsutil cp gs://bucketName /path/on/pod`



Google Cloud



MU-HPDI/Nautilus



University of Missouri



# Migrating Data to Nautilus: Using Nautilus S3 Storage



- ▶ NRP provides S3 bucket storage for free to Nautilus users upon request
- ▶ Advantages:
  - ▶ Free
  - ▶ High-throughput link to Nautilus cluster
  - ▶ Integration with S3-compatible software
  - ▶ Capability to handle large amounts of data
- ▶ Disadvantages:
  - ▶ Must request access
  - ▶ Setup of cloud integration
  - ▶ Staging of data into Nautilus S3
- ▶ How To:
  - ▶ Request access to cloud storage to receive Access ID and Keys
  - ▶ Install `rc1one` or similar tool at data source and copy data from source to Nautilus S3
  - ▶ Create `rc1one` or similar tool pod on Nautilus cluster and copy from S3 to Persistent Volume







**High-Performance Data-Intensive  
Computing Systems Laboratory**

Part 2-B

# Deploying Scikit-Learn Jobs to Nautilus for Machine Learning

Introduction to Containers and Kubernetes

IEEE International Conference on Big Data

December 2023



University of Missouri



# Prerequisites

- ▶ To run jobs on Nautilus, the following prerequisites must be met:
  - ▶ You have access to Nautilus and have been assigned a namespace
  - ▶ You have a container published to a public registry with the necessary code to perform the ML task
  - ▶ You have the data for the ML task on a persistent volume in the cluster



Access



ML  
Container



Data

# Access the MachineLearning\_K8s Notebook

Follow along activity



University of Missouri





**High-Performance Data-Intensive  
Computing Systems Laboratory**

# Part 2-C Deploying GPU Jobs to Nautilus for Computer Vision Applications

Introduction to Containers and Kubernetes

IEEE International Conference on Big Data

December 2023



University of Missouri



# Prerequisites - Review + GPU differences

- ▶ To run GPU jobs on Nautilus, the following prerequisites must be met:
  - ▶ You have access to Nautilus and have been assigned a namespace
  - ▶ You have a **GPU enabled** container published to a public registry with the necessary code to perform the CV task
  - ▶ You have the data for the CV task on a persistent volume in the cluster



Access



**GPU**

Container



Data



# Access the DeepLearning\_K8s Notebook

Follow along activity



University of Missouri





# High-Performance Data-Intensive Computing Systems Laboratory

## Part 2-D Automating Jobs on Nautilus using Bash and Python

Introduction to Containers and Kubernetes

IEEE International Conference on Big Data

December 2023



University of Missouri






# Deep Learning on Nautilus: Semantic segmentation

- ▶ Iterations of Training Completed: **515,550**
- ▶ Number of images Processed: **7,070,400**
- ▶ Trainable Parameters Optimized: **23 millions per model**
- ▶ The time it took to prepare the experimental set up and to run all the training sessions in parallel is **12 hours**
- ▶ The actual time it would have take to train is **21 days 12 hours 45 minutes**




# Deep Learning on Nautilus: Transformer Research




9 Deep Neural  
Architectures




27 Trained Models



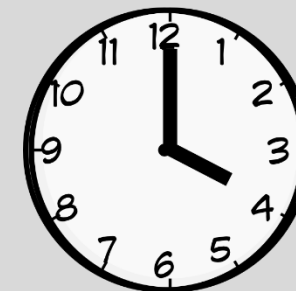
124.74 TB  
Imagery  
Processed



3 Open Source  
HR-RSI Datasets



- 8,100 Epochs
- 30M iterations
- 1.7B parameters



Wall Clock Time:  
76 days, 10 hours





# MU-HPDI/Nautilus - A resource for exploration

- ▶ Sample Dockerfiles
- ▶ Sample Kubernetes YAML File
- ▶ Wiki with detailed walkthroughs for:
  - ▶ Getting Started
  - ▶ Creating PVC
  - ▶ Creating Pods
  - ▶ Creating Jobs

Check back  
occasionally, we are  
adding more  
content and recipes  
for scaling

Machine Learning  
and Deep Learning

